

Kavya Agar

kavyaagar0@gmail.com | kavyaagar.com | github.com/Kavya-Agar | linkedin.com/in/kavya-agar

EDUCATION

Texas A&M University — College Station, TX

Aug 2024 – May 2027

B.S. in Computer Science | GPA: 3.5 | Dean's List | Minors: Statistics, Business

Coursework: Data Structures & Algorithms, Operating Systems, Computer Systems, Software Engineering, Computer Graphics, Database Management, Statistics

SKILLS

Languages: Java, Python, C/C++, JavaScript, TypeScript, HTML/CSS, Swift, SQL, R

Frameworks: React, Django, Flask, Selenium, REST APIs, PyTorch, TensorFlow, HuggingFace, LangChain, Scikit-learn, JUnit, Pandas

Cloud & DevOps: AWS (EC2, S3, SageMaker, Step Functions, Lambda), Docker, Kubernetes, Terraform, Git, CI/CD (Jenkins, GitHub Actions, ArgoCD)

Data & ML: PostgreSQL, MySQL, MongoDB, DynamoDB, Cassandra, Elasticsearch, NLP, Deep Learning, MLOps, ETL Pipelines

EXPERIENCE

JPMorgan Chase & Co. — Houston, TX

Jun 2023 – Aug 2023

Software Engineering Intern | Team Lead

- Engineered a multithreaded Java data pipeline for a distributed system migrating 10TB+ from Apache Cassandra to AWS S3, implementing custom Parquet conversion and improving migration efficiency by 90%
- Collaborated with 3 senior engineers through 10+ code reviews to harden system reliability, performance, and data integrity, writing unit tests and validating end-to-end pipeline correctness
- Led daily Agile stand-ups for an 8-person cross-functional team, coordinating sprint execution and presenting the final solution to Executive Directors

Aspen Technology — Houston, TX

Jan 2024 – Apr 2024

Technical Business Analyst Intern

- Built automated Python web-scraping pipelines using BeautifulSoup and Selenium to analyze 5+ competing analytics platforms, reducing manual research time by 60%
- Partnered with PMs and engineers to translate data insights into product strategy decisions, contributing to an estimated 15% revenue impact

PROJECTS

Financial NER Analyzer | *Python, PyTorch, HuggingFace, AWS (Step Functions, SageMaker, Textract), React, Docker*

- Built a production-ready document classifier supporting multiple financial document types with custom entity recognition for transaction amounts, dates, account numbers, and payment identifiers
- Fine-tuned DistilBERT on 116K+ tokens, achieving 98% accuracy and reducing manual document review time by 70%
- Architected a fault-tolerant AWS pipeline (Step Functions, Textract, SageMaker) processing 500+ documents/week with automatic error handling and retry logic

Stock Sentiment Dashboard | *Python, AWS (SageMaker, DynamoDB), HuggingFace, Terraform, Flask, React, Docker*

- Deployed an end-to-end ML pipeline using Terraform and Docker serving 50+ users with real-time sentiment analysis from Reddit and Finnhub
- Implemented ensemble NLP inference (HuggingFace Transformers + VADER) achieving 78% accuracy and configured CloudWatch alarms for pipeline health monitoring
- Designed a full-stack application with Flask and React to visualize sentiment trends, storing live data with DynamoDB

FIN-ance | *Django, React, TensorFlow, SwiftUI, PostgreSQL, Docker, OAuth 2.0*

- Built a cross-platform payment and expense-tracking app processing financial transactions with ML-based spending forecasts (90% accuracy), serving 100+ beta users
- Designed a Django REST API handling 1,000+ daily requests with OAuth 2.0 authentication and real-time mobile/web sync, ensuring secure transaction handling
- Implemented iOS client using SwiftUI with WidgetKit for iPhone home screen expense tracking

CERTIFICATIONS

AWS Certified Cloud Practitioner (CLF-C02)

IBM DevOps & Software Engineering

Databricks GenAI Fundamentals

BCG GenAI Job Simulation